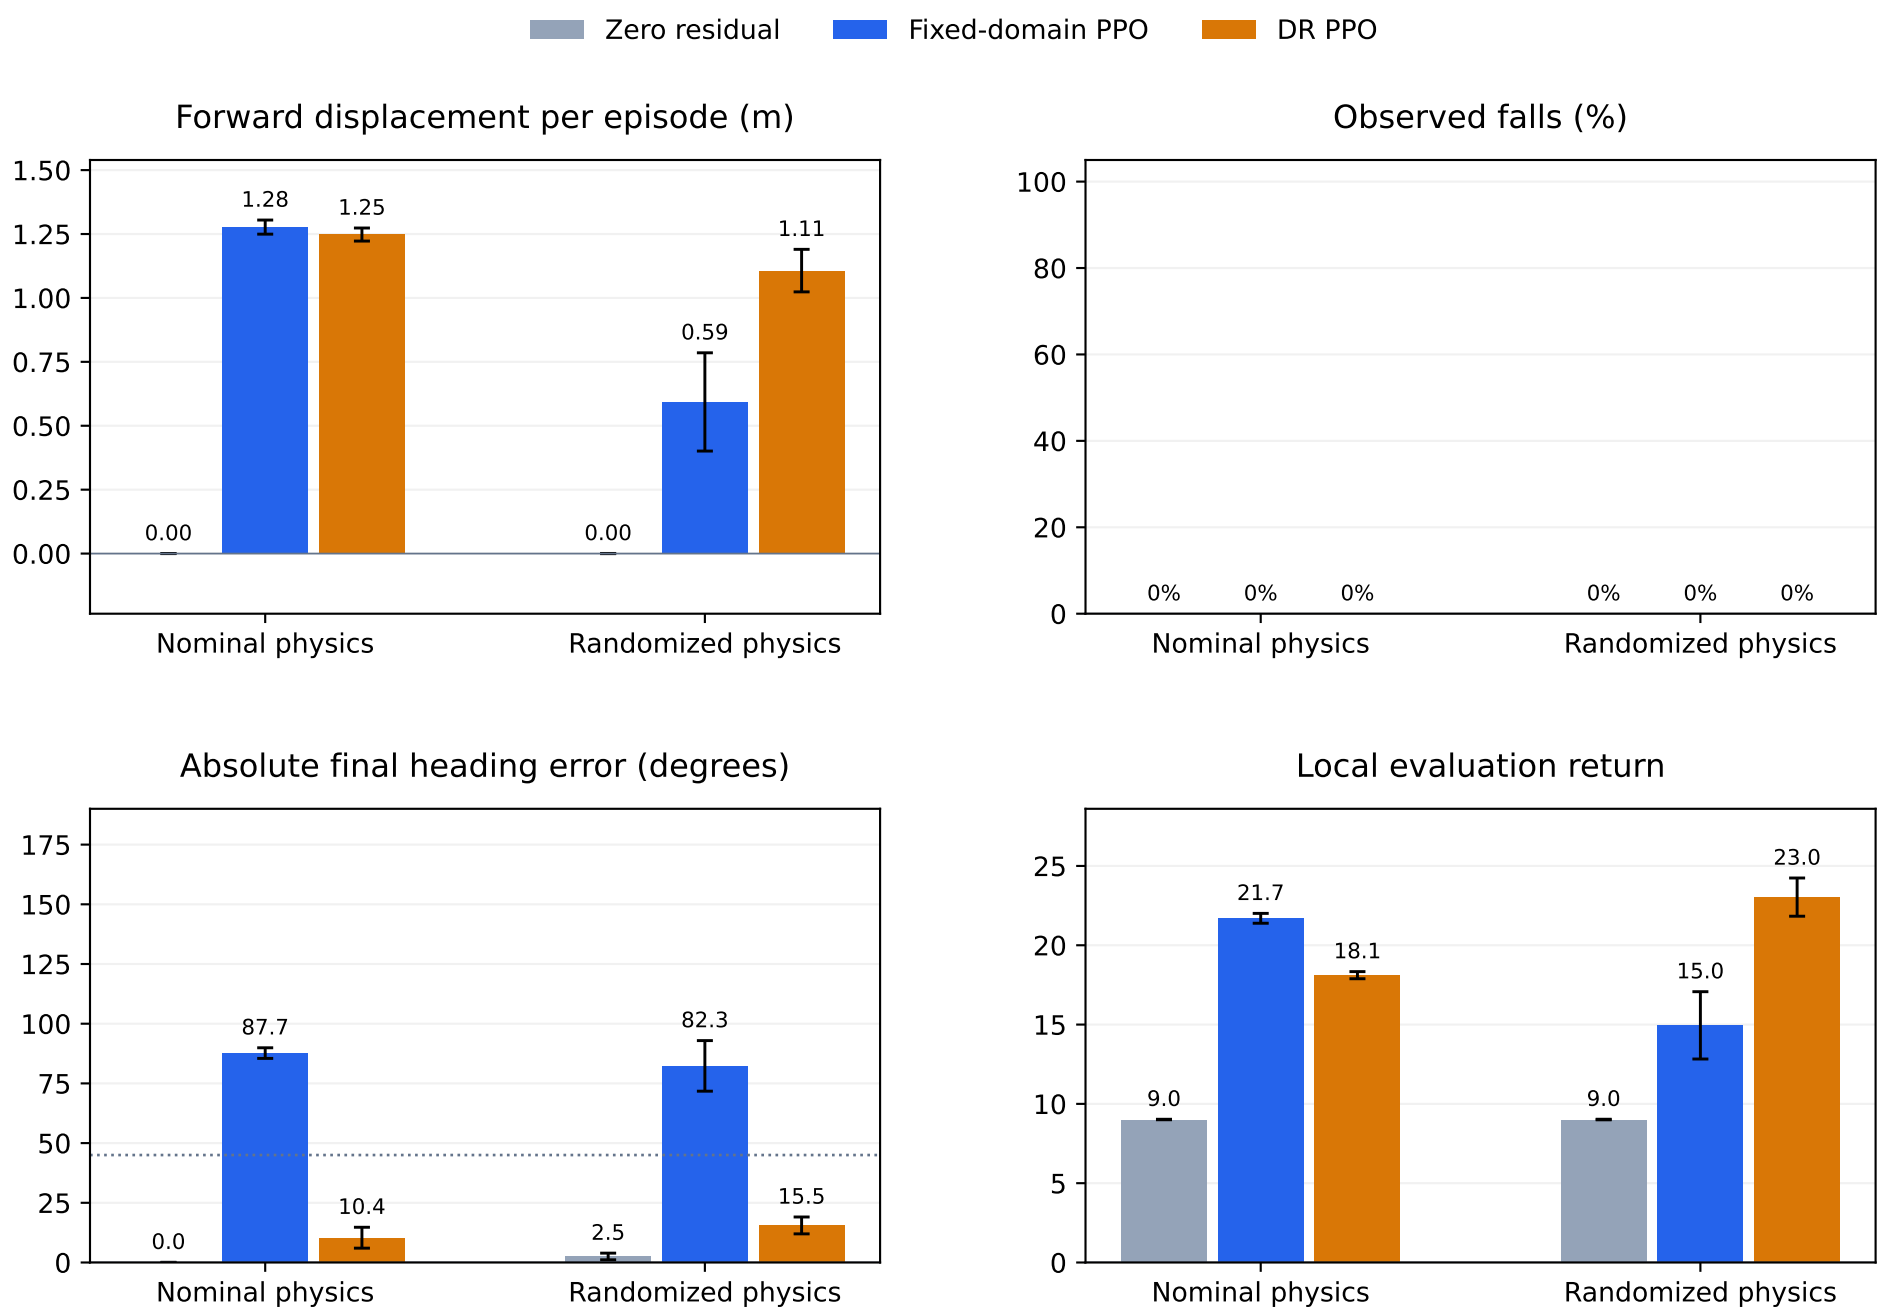
Sesame PPO | paired held-out evaluation | DR: 250 updates
Zero residual
Fixed-domain PPO
DR PPO
Forward displacement per episode (m)
1.50
1.25
1.00
0.75
0.50
0.25
0.00
1.28
1.25
1.11
0.59
0.00
0.00
Nominal physics
Randomized physics
Observed falls (%)
100
80
60
40
20
0
0%
0%
0%
0%
0%
0%
Nominal physics
Randomized physics
Absolute final heading error (degrees)
175
150
125
100
75
50
25
0
87.7
82.3
10.4
15.5
0.0
2.5
Nominal physics
Randomized physics
Local evaluation return
25
20
15
10
5
0
21.7
18.1
23.0
15.0
9.0
9.0
Nominal physics
Randomized physics
64 episodes per policy/domain; seed 2042. Bars: episode mean; error bars: 1 SD, not confidence intervals.
One training seed per learned policy. Heading is relative to world +x; reward was not changed to fix sideways motion.