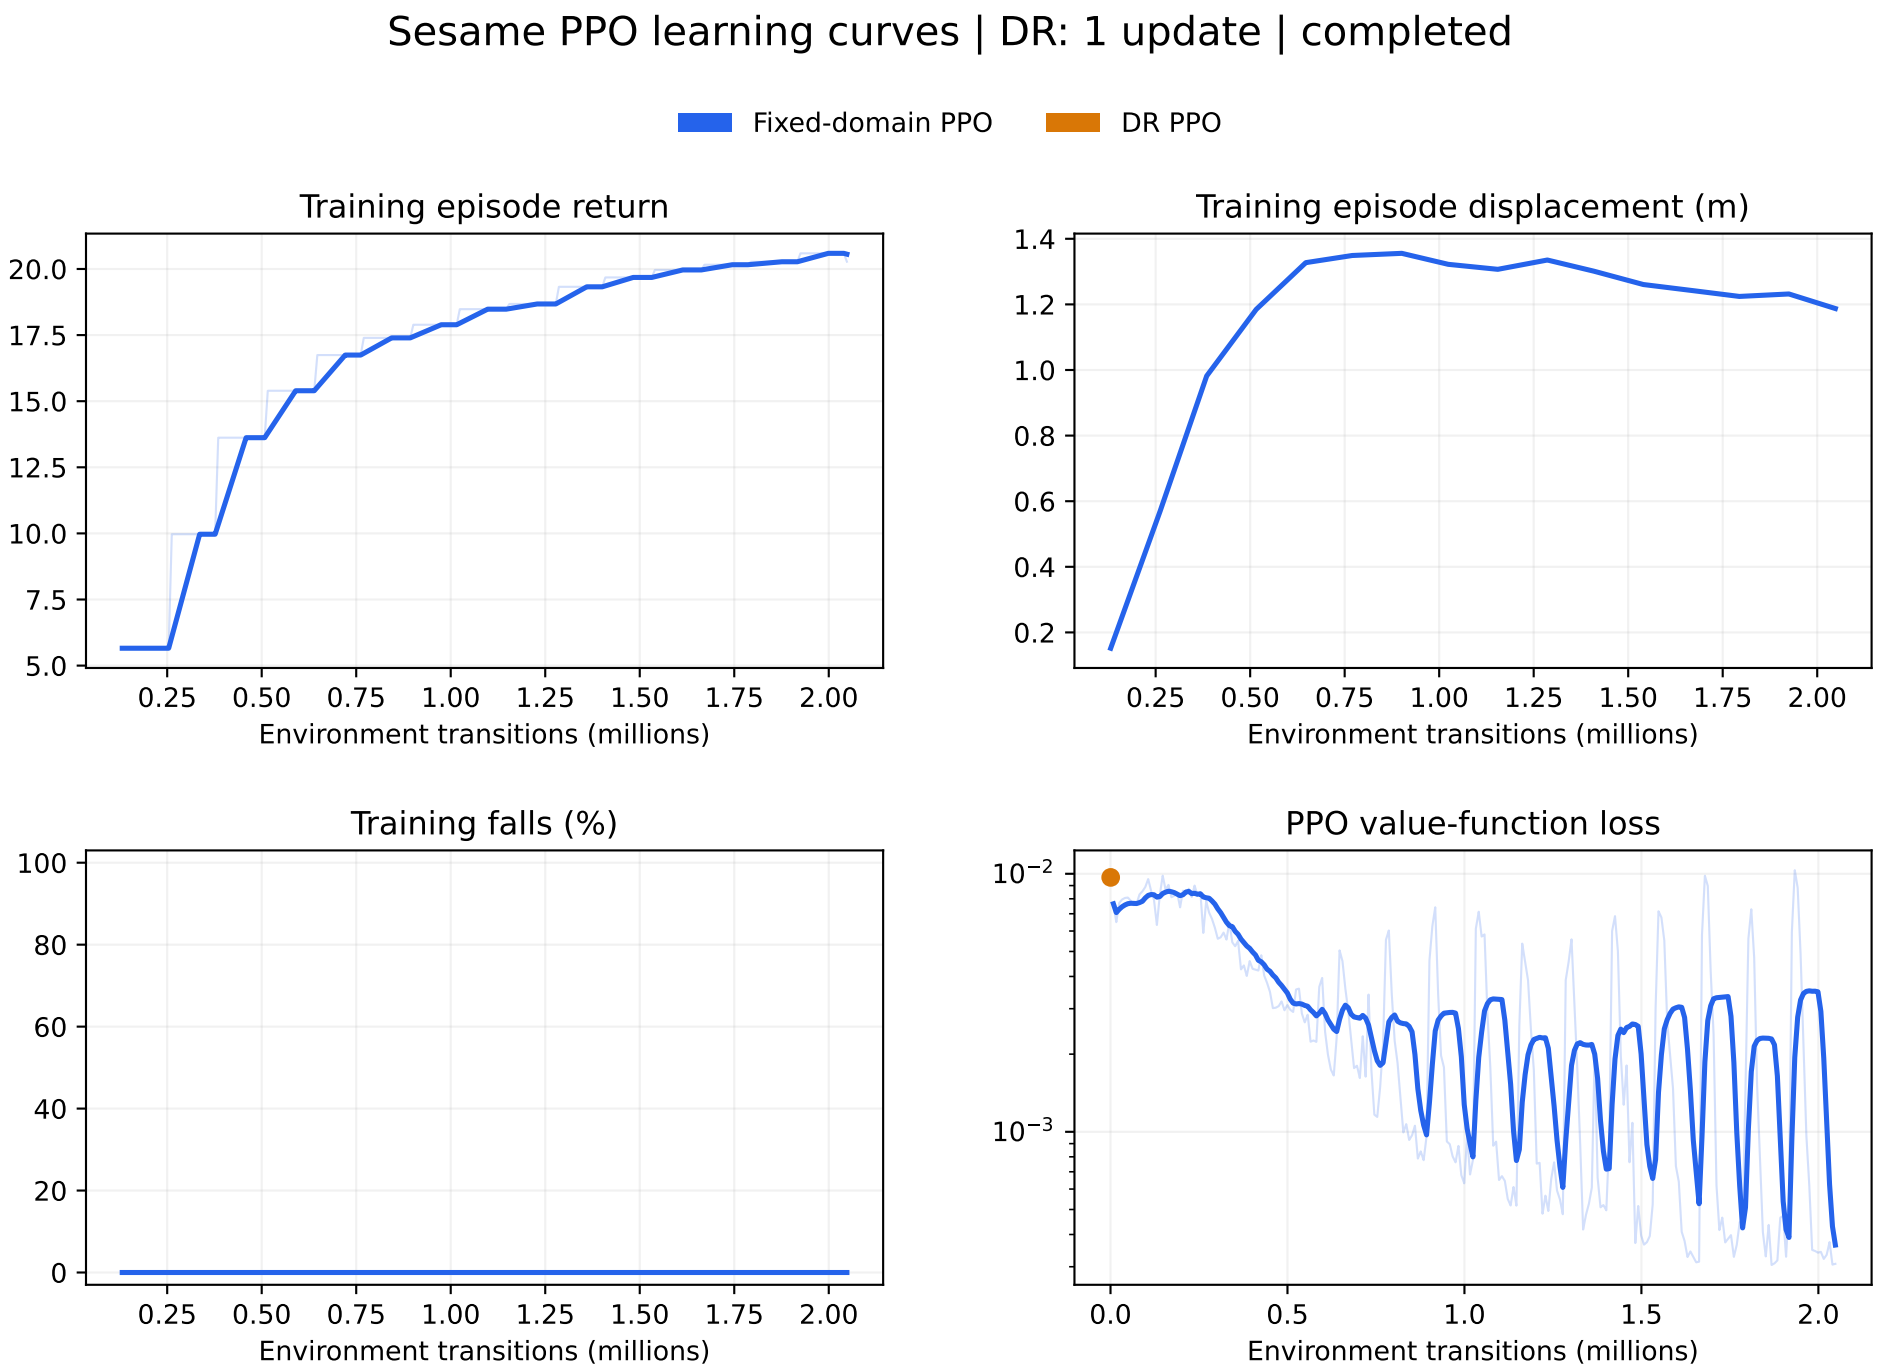
Sesame PPO learning curves | DR: 1 update | completed
Fixed-domain PPO
DR PPO
Training episode return
20.0
17.5
15.0
12.5
10.0
7.5
5.0
0.25
0.50
0.75
1.00
1.25
1.50
1.75
2.00
Environment transitions (millions)
Training episode displacement (m)
1.4
1.2
1.0
0.8
0.6
0.4
0.2
0.25
0.50
0.75
1.00
1.25
1.50
1.75
2.00
Environment transitions (millions)
Training falls (%)
100
80
60
40
20
0
0.25
0.50
0.75
1.00
1.25
1.50
1.75
2.00
Environment transitions (millions)
PPO value-function loss
10^{-2}
10^{-3}
0.0
0.5
1.0
1.5
2.0
Environment transitions (millions)
Raw values are faint; solid lines show a trailing 10-update mean. Training domains differ; use the paired evaluation for comparison.