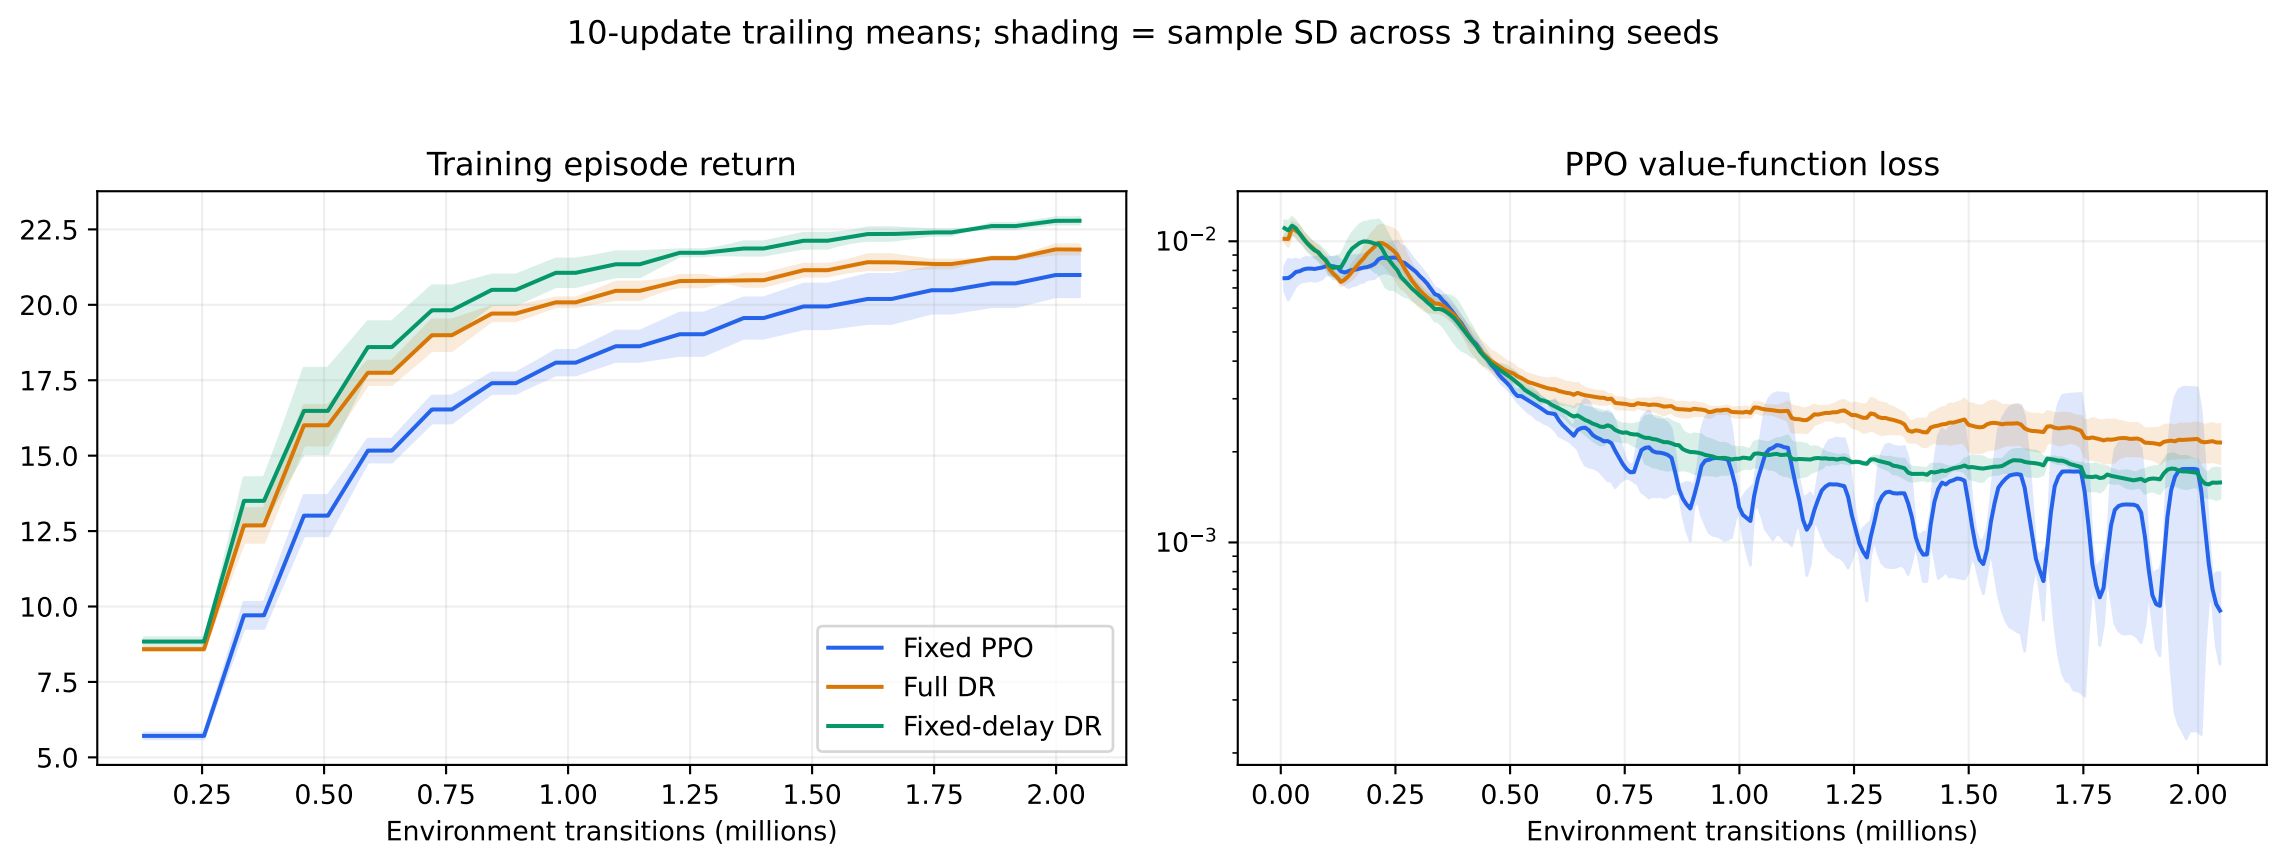
10-update trailing means; shading = sample SD across 3 training seeds
Training episode return
PPO value-function loss
Fixed PPO
Full DR
Fixed-delay DR
Environment transitions (millions)
Environment transitions (millions)
Training physical conditions differ across methods. Use paired held-out evaluation to compare performance.