Sesame PPO | 3 training seeds × 3 methods

Fixed PPO | Full DR | Fixed-delay DR

Forward speed tracking RMSE (m/s)

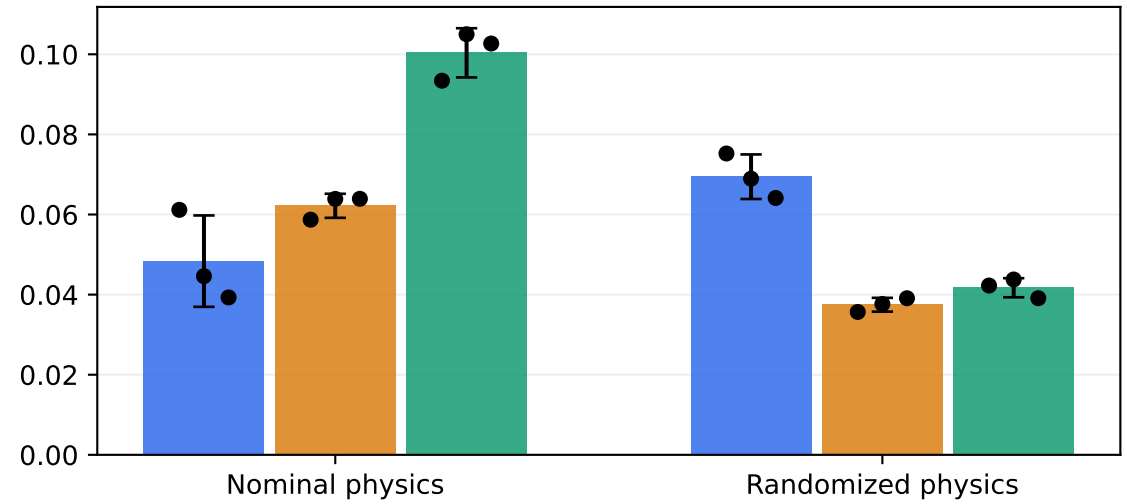

Forward displacement (m)

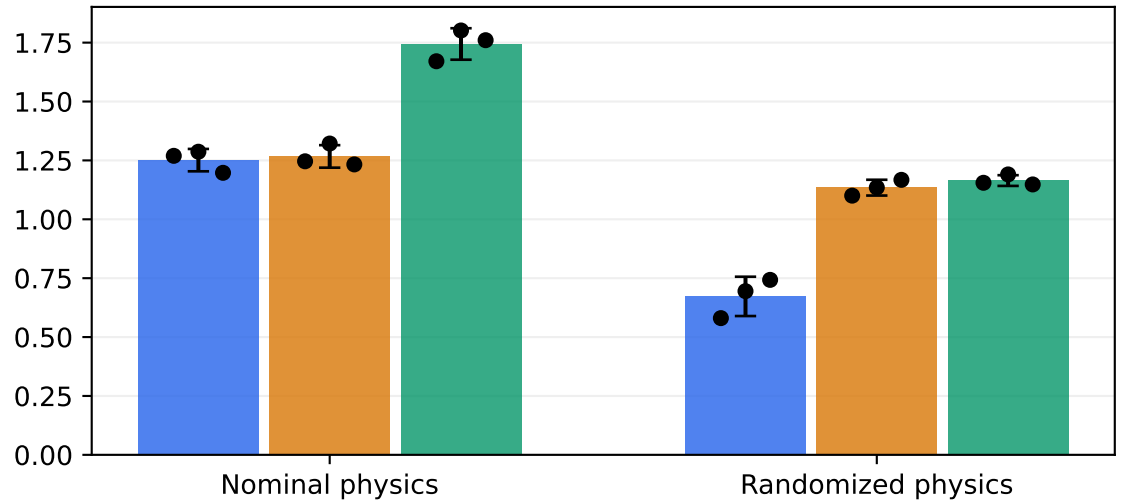

Absolute final heading (degrees)

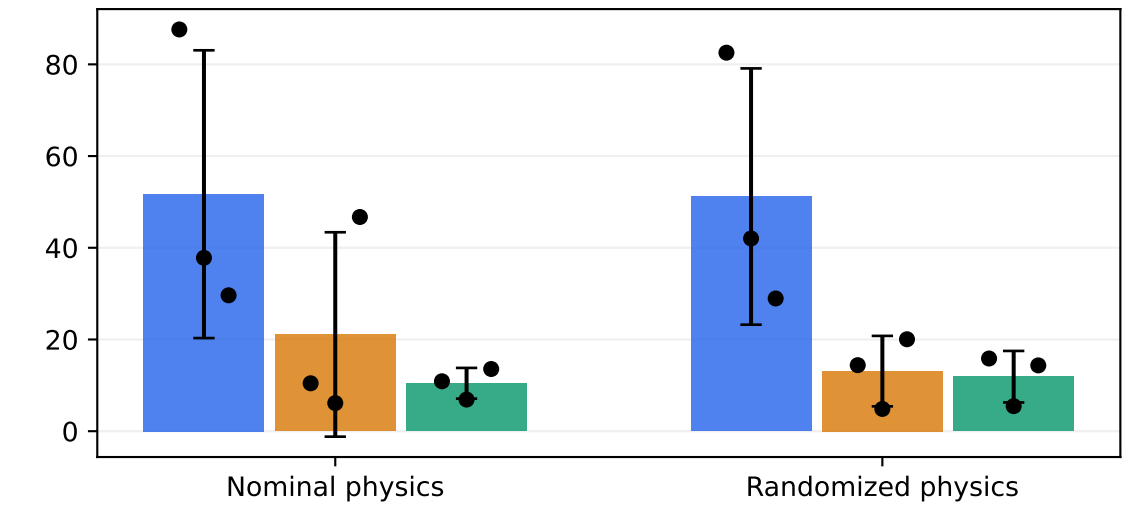

Absolute final lateral offset (m)

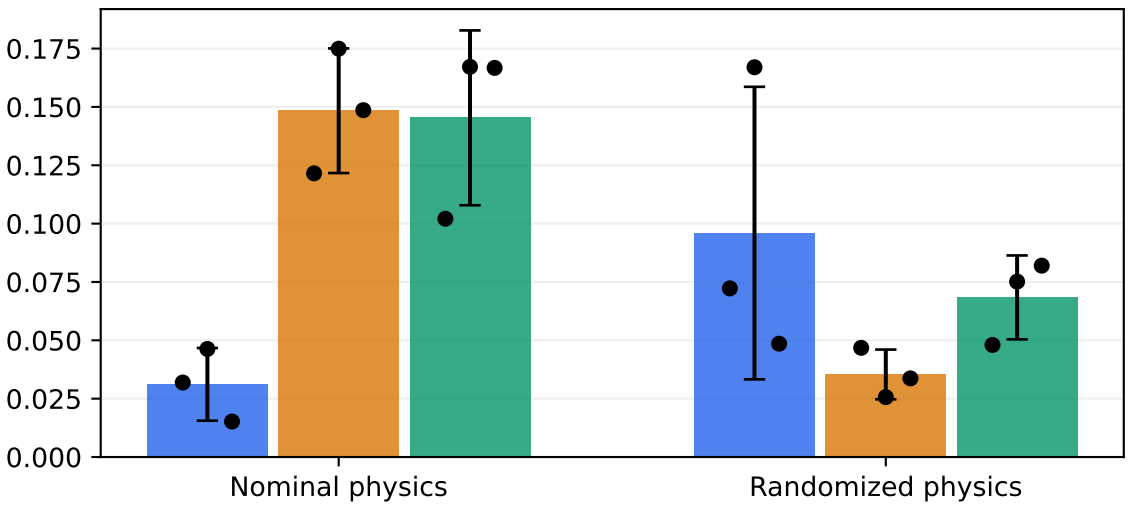

Unchanged local evaluation return

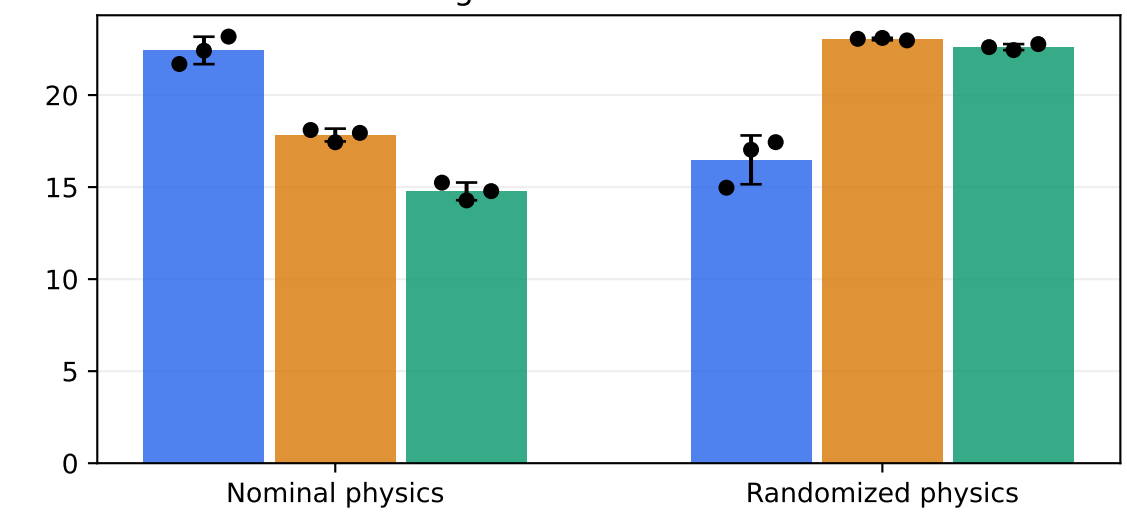

Observed fall fraction

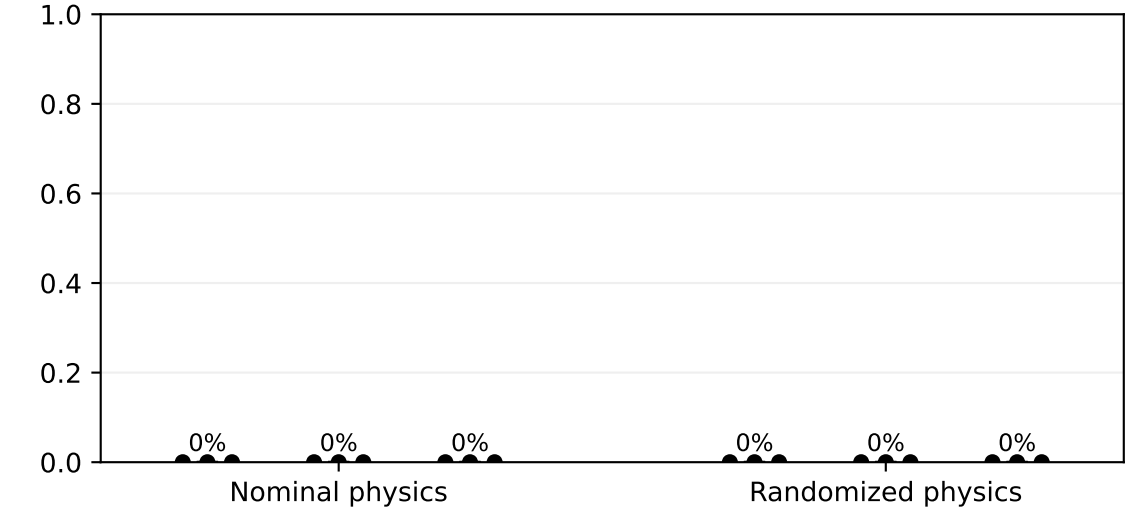

Dots: training-seed means (128 paired episodes each). Error bars: sample SD across 3 training seeds, not confidence intervals.
Same evaluation physics for every method; unchanged local reward. Target/joint derivatives sampled at 50 Hz.